



开Oracle调优鹰眼，深入理解 AWR性能报告



刘相兵(Maclean Liu)
liu.maclean@gmail.com

ORA-ALLSTARS Exadata用
户组QQ群:23549328

About Me

ORACLE®
Certified Master
Oracle Database 10g
Administrator

ORACLE®
Certified Master
Oracle Database 11g
Administrator

- Email: liu.maclean@gmail.com
- Blog: www.askmaclean.com
- Oracle Certified Database Administrator Master 10g and 11g
- Over 7 years experience with Oracle DBA technology
- Over 8 years experience with Linux technology
- Member Independent Oracle Users Group
- Member All China Users Group
- Presents for advanced Oracle topics: RAC, DataGuard, Performance Tuning and Oracle Internal.



For the Complete Technology & Database Professional



MEMBER

Hawk Eyes 看AWR的鹰眼= 基础理论 夯实+看过500份以上AWR



啥是AWR?

AWR (Automatic Workload Repository)

一堆历史性能数据，放在**SYSAUX**表空间上，**AWR**和**SYSAUX**都是**10g**出现的，是**Oracle**调优的关键特性；大约**1999**年左右开始开发，已经有**15**年历史

默认快照间隔**1**小时，**10g**保存**7**天、**11g**保存**8**天；可以通过
DBMS_WORKLOAD_REPOSITORY.MODIFY_SNAPSHOT_SETTINGS
修改

DBA_HIST_WR_CONTROL

AWR程序核心是**dbms_workload_repository**包

@?/rdbms/admin/awrrpt 本实例

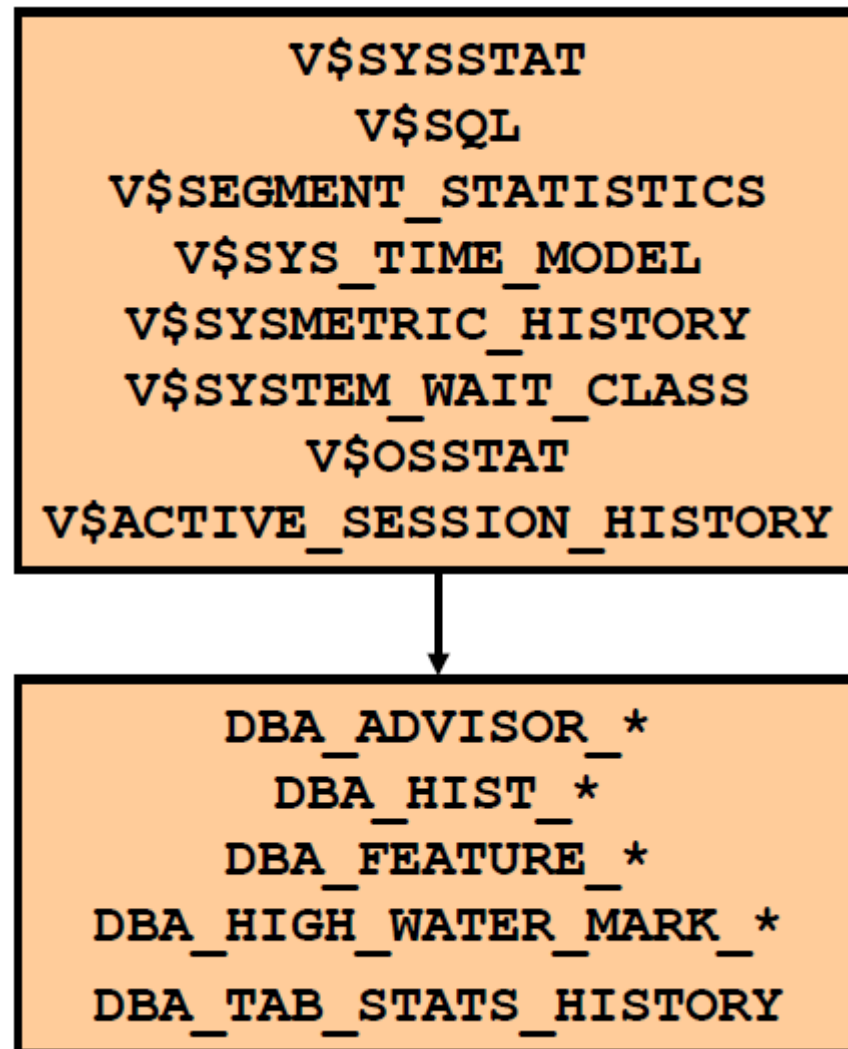
@?/rdbms/admin/awrrpti RAC中选择实例号

AWR数据从哪里来，放哪里去？

基础指标统计：

- ◆ SQL和优化器指标
- ◆ OS指标
- ◆ 等待事件类型
- ◆ 时间指标

- 度量
- **ASH Active Session History**
- 建议器advisor
- 快照指标
- 数据库特性使用情况

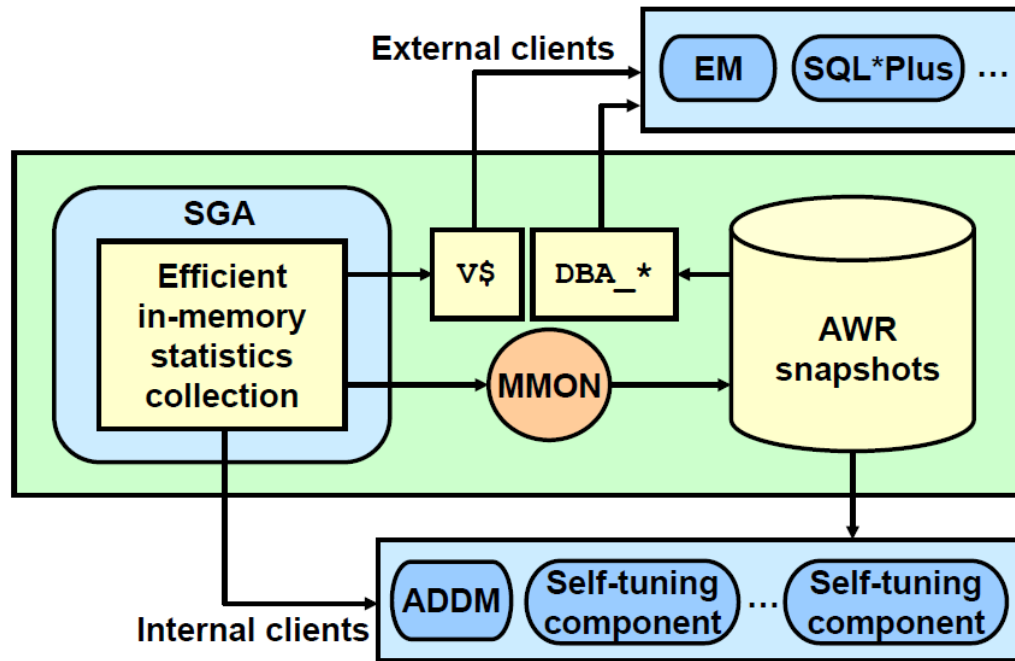


谁维护AWR?

主要是**MMON(Manageability Monitor Process)**和它的小工进程(m00x)

MMON的功能包括:

1. 启动**slave**进程m00x去做**AWR**快照
2. 当某个度量阈值被超过时发出**alert**告警
3. 为最近改变过的**SQL**对象捕获指标信息



mmon若hang则意味着AWR不可用!!

AWR小技巧

手动执行一个快照:

Exec dbms_workload_repository.create_snapshot; (这个要背出来哦, 用的时候去翻手册, 丢脸哦 ☺!)

创建一个AWR基线

Exec

DBMS_WORKLOAD_REPOSITORY.CREATE_BASELINE(start_snap_id, end_snap_id, baseline_name);

@?/rdbms/admin/awrddrpt AWR 比对报告

@?/rdbms/admin/awrgrpt RAC 全局AWR

自动生成AWR HTML报告:

http://www.oracle-base.com/dba/10g/generate_multiple_awr_reports.sql

AWR开篇 卷首灌顶，惊鸿一瞥

祇園精舎の鐘の聲

諸行無常の響あり

娑羅雙樹の色

盛者必衰の理あり

おぼなる人も久からず

唯春の夜の夢の如く

たけき者も遂にはほろびぬ

偏に風の前の塵に同じ

性能无
性常

盛者不
盛凋必
者优衰

DB TIME

WORKLOAD REPOSITORY report for

| DB Name | DB Id | Instance | Inst num | Startup Time | Release | RAC |
|---------|-----------|----------|----------|-----------------|------------|-----|
| MACDB | 381251694 | MACDB1 | 1 | 12-Dec-11 14:12 | 11.2.0.1.0 | YES |

版本和是否为RAC

| Host Name | Platform | CPUs | Cores | Sockets | Memory (GB) |
|----------------------|------------------|------|-------|---------|-------------|
| MACDB01.zeradata.com | Linux x86 64-bit | 16 | 8 | 2 | 70.60 |

平台、CPU和物理内存

| | Snap Id | Snap Time | Sessions | Cursors/Session |
|-------------|---------|--------------------|----------|-----------------|
| Begin Snap: | 17323 | 17-Feb-12 21:00:16 | 106 | 3.9 |
| End Snap: | 17324 | 17-Feb-12 22:00:09 | 103 | 3.9 |
| Elapsed: | | 59.88 (mins) | | |
| DB Time: | | 42.74 (mins) | | |

session总数和平均每个session打开多少cursor

逝去时间和DB TIME(超重要性能指标哦!), 逝去时间视乎需求可长可短!

信息量很巨大哦!!

Elapsed快照逝去时间, 如果为了诊断特定时段性能问题则不宜过长**15分钟~2、3个小时**。如果是看全天负载那么可以长一些。

最常见是**60分钟**后者**120分钟**, 视乎实际需求, 无成法。

Cursors/session → **open_cursors** **ORA-100**诊断

DB TIME

DB TIME= 所有前台**session**花费在**database**调用上的总和时间

- 注意是前台进程**foreground sessions**
- 包括**CPU**时间、**IO Time**、和其他一系列**非空闲**等待时间，别忘了**cpu on queue time**

DB TIME 不等于 响应时间，**DB TIME**高了未必响应慢，**DB TIME**低了未必响应快

DB Time描绘了数据库总体负载，但要和**elapsed time**逝去时间结合起来。

Average Active Session AAS= DB time/Elapsed Time

DB Time =60 min , Elapsed Time =60 min AAS=60/60=1 负载一般

DB Time= 1min , Elapsed Time= 60 min AAS= 1/60 负载很轻

DB Time= 60000 min, Elapsed Time= 60 min AAS=1000 → 系统hang了吧？

DB TIME= DB CPU + NON-Idle Wait + Wait on CPU queue

DB TIME= DB CPU + Non-Idle Wait + Wait on CPU queue

如果仅有2个逻辑CPU，而2个session在60分钟都没等待事件，一直跑在CPU上，那么：

DB CPU= 2 * 60 mins ， DB Time = 2* 60 + 0 + 0 =120

AAS = 120/60=2 正好等于OS load 2。

如果有3个session都100%仅消耗CPU，那么总有一个要wait on queue

DB CPU = 2* 60 mins ， wait on CPU queue= 60 mins


AAS= (120+ 60)/60=3 → 主机load 亦为3，此时vmstat 看waiting for run time

真实世界中？ DB Cpu = xx mins ， Non-Idle Wait= enq:TX + cursor pin S on X + latch : xxx + db file sequential read + 阿猫阿狗

查最近7天的DB Time

```
WITH sysstat AS
(select sn.begin_interval_time begin_interval_time,
      sn.end_interval_time end_interval_time,
      ss.stat_name stat_name,
      ss.value e_value,
      lag(ss.value, 1) over(order by ss.snap_id) b_value
 from dba_hist_sysstat ss, dba_hist_snapshot sn
 where trunc(sn.begin_interval_time) >= sysdate - 7
    and ss.snap_id = sn.snap_id
    and ss.dbid = sn.dbid
    and ss.instance_number = sn.instance_number
    and ss.dbid = (select dbid from v$database)
    and ss.instance_number = (select instance_number from v$instance)
    and ss.stat_name = 'DB time')
select to_char (BEGIN_INTERVAL_TIME, 'mm-dd hh24:mi') || to_char (END_INTERVAL_TIME, '
hh24:mi') date_time, stat_name, round((e_value - nvl(b_value, 0)) / (extract(day
from(end_interval_time - begin_interval_time)) * 24 * 60 * 60 + extract(hour
from(end_interval_time - begin_interval_time)) * 60 * 60 + extract(minute
from(end_interval_time - begin_interval_time)) * 60 + extract(second
from(end_interval_time - begin_interval_time))), 0) per_sec
 from sysstat
 where(e_value - nvl(b_value, 0)) > 0 and nvl(b_value, 0) > 0
/
```

DB Time折线图：文似看山不喜平



实战DB Time案例

同一套库 1月 VS 2月

| | Snap Id | Snap Time | Sessions | Cursors/Session |
|-------------|---------|--------------------|----------|-----------------|
| Begin Snap: | 16193 | 01-Jan-12 21:00:04 | 237 | 182.1 |
| End Snap: | 16194 | 01-Jan-12 22:00:11 | 237 | 182.1 |
| Elapsed: | | 60.12 (mins) | | |
| DB Time: | | 796.28 (mins) | | |


| | Snap Id | Snap Time | Sessions | Cursors/Session |
|-------------|---------|--------------------|----------|-----------------|
| Begin Snap: | 17323 | 17-Feb-12 21:00:16 | 245 | 176.4 |
| End Snap: | 17324 | 17-Feb-12 22:00:09 | 245 | 176.3 |
| Elapsed: | | 59.88 (mins) | | |
| DB Time: | | 3,094.52 (mins) | | |

| Statistic Name | Time (s) | % of DB Time |
|--|-----------|--------------|
| sql execute elapsed time | 26,173.49 | 54.78 |
| DB CPU | 14,702.54 | 30.77 |
| connection management call elapsed time | 300.54 | 0.63 |
| PL/SQL execution elapsed time | 45.11 | 0.09 |
| repeated bind elapsed time | 18.84 | 0.04 |
| parse time elapsed | 17.40 | 0.04 |
| hard parse elapsed time | 0.84 | 0.00 |
| hard parse (sharing criteria) elapsed time | 0.64 | 0.00 |
| sequence load elapsed time | 0.47 | 0.00 |
| DB time | 47,776.99 | |
| background elapsed time | 11,093.26 | |
| background cpu time | 1,409.87 | |

| Statistic Name | Time (s) | % of DB Time |
|--|------------|--------------|
| sql execute elapsed time | 165,268.77 | 89.01 |
| DB CPU | 10,273.18 | 5.53 |
| connection management call elapsed time | 27.56 | 0.01 |
| PL/SQL execution elapsed time | 14.50 | 0.01 |
| repeated bind elapsed time | 14.39 | 0.01 |
| parse time elapsed | 3.00 | 0.00 |
| hard parse elapsed time | 0.61 | 0.00 |
| hard parse (sharing criteria) elapsed time | 0.44 | 0.00 |
| sequence load elapsed time | 0.25 | 0.00 |
| DB time | 185,670.96 | |
| background elapsed time | 11,804.35 | |
| background cpu time | 1,060.59 | |

AAS平均活动会话==与ASH结合

At time t we have 2 active sessions



| | Sample Time | Data Source |
|------------------------------|--------------------|---------------------------|
| Analysis Begin Time: | 01-Feb-13 00:11:55 | V\$ACTIVE_SESSION_HISTORY |
| Analysis End Time: | 01-Feb-13 01:11:56 | V\$ACTIVE_SESSION_HISTORY |
| Elapsed Time: | 60.0 (mins) | |
| Sample Count: | 666,824 | |
| Average Active Sessions: | 185.18 | |
| Avg. Active Session per CPU: | 1.45 | |
| Report Target: | None specified | |

ASH报告开
篇即为AAS

Cache Size 兵马未动，缓存先行

内存管理方式: MSMM、ASMM(sga_target)、AMM(memory_target)

小内存有小内存的问题，大内存有大内存的麻烦！ORA-04031????!

Buffer cache和shared pool size的 begin/end值在ASMM、AMM和11gR2 MSMM下可是会动的哦！

这里说 shared pool一直收缩，则在shrink过程中一些row cache 对象被lock住可能导致前台row cache lock等解析等待，最后别让shared pool shrink。

如果这里shared pool一直在grow，那说明shared pool原有大小不足以满足需求(可能是大量硬解析)，结合下文的解析信息和SGA breakdown来一起诊断问题。

Cache Sizes

| | Begin | End | | |
|-------------------|---------|---------|-----------------|---------|
| Buffer Cache: | 23,040M | 23,040M | Std Block Size: | 8K |
| Shared Pool Size: | 5,888M | 5,888M | Log Buffer: | 23,568K |

Load Profile

| | Per Second | Per Transaction | Per Exec | Per Call |
|------------------|--------------|-----------------|----------|----------|
| DB Time(s): | 152.8 | 0.1 | 0.01 | 0.00 |
| DB CPU(s): | 12.0 | 0.0 | 0.00 | 0.00 |
| Redo size: | 25,188,867.2 | 13,533.4 | | |
| Logical reads: | 707,596.0 | 380.2 | | |
| Block changes: | 59,059.6 | 31.7 | | |
| Physical reads: | 12,011.1 | 6.5 | | |
| Physical writes: | 5,152.5 | 2.8 | | |
| User calls: | 30,638.1 | 16.5 | | |
| Parses: | 1,492.7 | 0.8 | | |
| Hard parses: | 491.9 | 0.3 | | |

SGA breakdown difference

- ordered by Pool, Name
- N/A value for Begin MB or End MB indicates the size of that Pool

| Pool | Name | Begin MB | End MB | % Diff |
|--------|---------------------------|----------|----------|--------|
| java | free memory | 1,792.00 | 1,792.00 | 0.00 |
| large | free memory | 5,112.19 | 5,112.19 | 0.00 |
| shared | ASH buffers | 254.00 | 254.00 | 0.00 |
| shared | Checkpoint queue | 1,024.06 | 1,024.06 | 0.00 |
| shared | FileOpenBlock | 248.32 | 248.32 | 0.00 |
| shared | KGH: NO ACCESS | 196.88 | 177.19 | -10.00 |
| shared | KGLH0 | 2,185.47 | 2,414.51 | 10.48 |
| shared | KGLHD | 392.50 | 464.37 | 18.31 |
| shared | KKSSP | 2,236.80 | 2,274.96 | 1.71 |
| shared | SQLA | 2,321.22 | 2,318.35 | -0.12 |
| shared | db_block_hash_buckets | 356.00 | 356.00 | 0.00 |
| shared | dbktb: trace buffer | 265.63 | 265.63 | 0.00 |
| shared | event statistics per sess | 319.53 | 319.53 | 0.00 |
| shared | free memory | 542.09 | 236.15 | -56.44 |
| shared | gcs resources | 1,072.24 | 1,072.24 | 0.00 |
| shared | gcs shadows | 625.47 | 625.47 | 0.00 |
| shared | ges big msg buffers | 248.31 | 248.31 | 0.00 |
| shared | ges enqueues | 377.05 | 377.05 | 0.00 |
| shared | ges resource | 263.63 | 263.63 | 0.00 |
| shared | kglsim object batch | 181.12 | 181.12 | 0.00 |
| shared | ksunfr: SSO free list | 297.85 | 297.85 | 0.00 |

Load Profile

Load Profile

| | Per Second | Per Transaction | Per Exec | Per Call |
|-------------------|-------------|-----------------|----------|----------|
| DB Time(s): | 36.5 | 0.2 | 0.01 | 0.00 |
| DB CPU(s): | 14.0 | 0.1 | 0.00 | 0.00 |
| Redo size: | 1,556,594.3 | 6,746.1 | | |
| Logical reads: | 1,579,406.7 | 6,845.0 | | |
| Block changes: | 7,569.8 | 32.8 | | |
| Physical reads: | 5,557.8 | 24.1 | | |
| Physical writes: | 385.6 | 1.7 | | |
| User calls: | 8,270.0 | 35.8 | | |
| Parses: | 1,744.9 | 7.6 | | |
| Hard parses: | 32.9 | 0.1 | | |
| W/A MB processed: | 4.7 | 0.0 | | |
| Logons: | 7.9 | 0.0 | | |
| Executes: | 3,884.3 | 16.8 | | |
| Rollbacks: | 1.1 | 0.0 | | |
| Transactions: | 230.7 | | | |

信息量太大！！

Load Profile

Redo size 单位 bytes，redo size可以用来估量update/insert/delete的频率，大的redo size往往对lgwr写日志，和arch归档造成I/O压力， Per Transaction可以用来分辨是 大量小事务， 还是少量大事务
如上例每秒redo 约1.5MB， 每个事务6k， 符合OLTP特征

Logical Read单位 次数*块数， 相当于“人*次”， 如上例 $1579406 * db_block_size = 12GB/s$ ， 逻辑读耗CPU， 主频和CPU核数都很重要， 逻辑读高则DB CPU往往高， 也往往可以看到latch: cache buffer chains等待。 大量OLTP系统(例如siebel)可以高达几十乃至上百Gbytes。

Block changes 单位 次数*块数， 描绘数据变化频率

Physical Read 单位次数*块数， 如上例 $5557 * 8k = 43MB/s$ ， 物理读消耗IO读， 体现在IOPS和吞吐量等不同纬度上；但减少物理读可能意味着消耗更多CPU。好的存储 每秒物理读能力达到几GB， 例如Exadata。

Load Profile

Physical writes 单位 次数*块数，主要是DBWR写datafile，也有direct path write。dbwr长期写出慢会导致定期log file switch(checkpoint not complete) 检查点无法完成的前台等待。

User Calls 单位次数，用户调用数，more details from internal

Parses，解析次数，包括软解析+硬解析，软解析优化得不好，则夸张地说几乎等于每秒SQL执行次数。即执行解析比1:1，而我们希望的是解析一次 到处运行哦！

Hard Parses：万恶之源。 **Cursor pin s on X**，
library cache: mutex X， **latch: row cache objects**
/shared pool.....。 硬解析最好少于每秒20次

Hard Parses实战案例

| | Per Second | Per Transaction | Per Exec | Per Call |
|-------------------|--------------|-----------------|----------|----------|
| DB Time(s): | 152.8 | 0.1 | 0.01 | 0.00 |
| DB CPU(s): | 12.0 | 0.0 | 0.00 | 0.00 |
| Redo size: | 25,188,867.2 | 13,533.4 | | |
| Logical reads: | 707,596.0 | 380.2 | | |
| Block changes: | 59,059.6 | 31.7 | | |
| Physical reads: | 12,011.1 | 6.5 | | |
| Physical writes: | 5,152.5 | 2.8 | | |
| User calls: | 30,638.1 | 16.5 | | |
| Parses: | 1,492.7 | 0.8 | | |
| Hard parses: | 491.9 | 0.3 | | |
| W/A MB processed: | 26.1 | 0.0 | | |
| Logons: | 0.9 | 0.0 | | |
| Executes: | 15,591.8 | 8.4 | | |
| Rollbacks: | 100.5 | 0.1 | | |
| Transactions: | 1,861.2 | | | |

Time Model Statistics

- Total time in database user-calls (DB Time): 477298s
- Statistics including the word "background" measure background process
- Ordered by % or DB time desc, Statistic name

| Statistic Name | Time (s) | % of DB Time |
|--|------------|--------------|
| sql execute elapsed time | 209,646.77 | 43.92 |
| parse time elapsed | 114,545.78 | 24.00 |
| DB CPU | 37,334.63 | 7.82 |
| hard parse elapsed time | 14,855.21 | 3.11 |
| PL/SQL execution elapsed time | 1,185.74 | 0.25 |
| hard parse (sharing criteria) elapsed time | 1,159.15 | 0.24 |
| hard parse (bind mismatch) elapsed time | 1,124.97 | 0.24 |
| connection management call elapsed time | 70.40 | 0.01 |
| sequence load elapsed time | 53.45 | 0.01 |
| PL/SQL compilation elapsed time | 32.57 | 0.01 |
| failed parse elapsed time | 6.63 | 0.00 |
| repeated bind elapsed time | 4.71 | 0.00 |
| DB time | 477,298.05 | |
| background elapsed time | 14,497.83 | |
| background cpu time | 3,389.66 | |

硬解析数和 *hard parse elapsed time* 对应，看一眼 *Time Model Statistics*，即可知该系统是否是解析敏感的

Hard Parses 实战案例

Top 5 Timed Foreground Events

| Event | Waits | Time(s) | Avg wait (ms) | % DB time | Wait Class |
|-----------------------------|------------|---------|---------------|-----------|-------------|
| log file sync | 5,500,239 | 137,224 | 25 | 28.75 | Commit |
| latch: row cache objects | 10,090,406 | 106,969 | 11 | 22.41 | Concurrency |
| db file sequential read | 9,067,793 | 65,018 | 7 | 13.62 | User I/O |
| latch: cache buffers chains | 7,500,108 | 38,049 | 5 | 7.97 | Concurrency |
| DB CPU | | 37,335 | | 7.82 | |

Latch:row cache objects 保护 row cache state object, 例如 dc_segments、dc_users

| | | | | |
|-------------------|---------------------------|---|-----------|-----------|
| redo writing | kcrfrsn | 0 | 2 | 15 |
| row cache objects | kqreqd: reget | 0 | 3,975,485 | 1,820,751 |
| row cache objects | kqrpre: find obj | 0 | 3,796,584 | 3,907,512 |
| row cache objects | kqrso | 0 | 1,412,501 | 2,324,276 |
| row cache objects | kqreqd | 0 | 1,152,438 | 2,283,575 |
| row cache objects | kqrbip | 0 | 941 | 161 |
| row cache objects | kqrigt | 0 | 872 | 434 |
| row cache objects | kqrbgl | 0 | 859 | 611 |
| row cache objects | kqrigt2 | 0 | 397 | 582 |
| row cache objects | kqrbtm: pop parent | 0 | 340 | 531 |
| row cache objects | kqrpre: init complete | 0 | 118 | 919 |
| row cache objects | ksucallcbksafely: kqrhngc | 0 | 106 | 0 |
| row cache objects | kqrbfr | 0 | 49 | 213 |
| row cache objects | kqrbpr: KQRRSNRL | 0 | 34 | 388 |

kqrpo parent cache object header
kqrso subordinate cache header
kqrpre() interface to get a row cache object

Hard Parses实战案例

Dictionary Cache Stats

- "Pct Misses" should be very low (< 2% in most cases)
- "Final Usage" is the number of cache entries being used

| Cache | Get Requests | Pct Miss | Scan Reqs | Pct Miss | Mod Reqs | Final Usage |
|----------------------|--------------|----------|-----------|----------|----------|-------------|
| dc_awr_control | 124 | 0.81 | 0 | | 2 | 1 |
| dc_constraints | 1,700 | 67.76 | 0 | | 1,700 | 0 |
| dc_files | 739,140 | 0.08 | 0 | | 0 | 762 |
| dc_global_oids | 1,034 | 3.87 | 0 | | 0 | 25 |
| dc_histogram_data | 26,206,700 | 0.32 | 0 | | 0 | 40,642 |
| dc_histogram_defs | 9,596,930 | 1.10 | 0 | | 156 | 37,960 |
| dc_object_grants | 21,627 | 0.49 | 0 | | 0 | 124 |
| dc_objects | 13,266,463 | 0.09 | 0 | | 1,592 | 17,297 |
| dc_profiles | 2,509 | 8.65 | 0 | | 0 | 0 |
| dc_rollback_segments | 250,813 | 0.00 | 0 | | 0 | 3,346 |
| dc_segments | 7,612,594 | 4.75 | 0 | | 3,768 | 600 |
| dc_sequences | 13,113 | 0.84 | 0 | | 13,113 | 24 |
| dc_table_scns | 75 | 100.00 | 0 | | 0 | 0 |
| dc_tablespace | 488,569 | 0.00 | 0 | | 0 | 49 |
| dc_users | 309,401,329 | 0.00 | 0 | | 0 | 134 |

Dc_histogram , 字典缓存row cache中的直方图信息

大量硬解析=>
直方图字典缓存争用

思路：解决硬解析，或者减少
*dc_histogram*争用

W/A MB processed

W/A MB processed : 单位MB W/A workarea workarea中处理的数据数量
结合 In-memory Sort%, sorts (disk) PGA Aggr一起看

Instance Efficiency Percentages (Target 100%)

| | | | |
|-------------------------------|-------|-------------------|--------|
| Buffer Nowait %: | 99.99 | Redo NoWait %: | 99.98 |
| Buffer Hit %: | 98.46 | In-memory Sort %: | 100.00 |
| Library Hit %: | 95.54 | Soft Parse %: | 67.05 |
| Execute to Parse %: | 90.43 | Latch Hit %: | 91.46 |
| Parse CPU to Parse Elapsed %: | 2.34 | % Non-Parse CPU: | 93.62 |

| | | | |
|------------------|-------------|------------|--------|
| sorts (disk) | 45 | 0.01 | 0.00 |
| sorts (memory) | 1,154,835 | 369.76 | 0.20 |
| sorts (rows) | 911,885,722 | 291,973.66 | 156.87 |
| sql area evicted | 1,489,921 | 477.05 | 0.26 |

W/A MB processed

W/A MB processed : 单位MB W/A workarea workarea中处理的数据数量
结合 In-memory Sort%, sorts (disk) PGA Aggr一起看

PGA Aggr Summary

- PGA cache hit % - percentage of W/A (WorkArea) data processed only in-memory

| PGA Cache Hit % | W/A MB Processed | Extra W/A MB Read/Written |
|-----------------|------------------|---------------------------|
| 64.24 | 96,871 | 53,918 |

[Back to Advisory Statistics](#)

[Back to Top](#)

PGA Aggr Target Stats

- B: Begin Snap E: End Snap (rows identified with B or E contain data which is absolute i.e. not diffed over the interval)
- Auto PGA Target - actual workarea memory target
- W/A PGA Used - amount of memory used for all Workareas (manual + auto)
- %PGA W/A Mem - percentage of PGA memory allocated to workareas
- %Auto W/A Mem - percentage of workarea memory controlled by Auto Mem Mgmt
- %Man W/A Mem - percentage of workarea memory under manual control

| | PGA Auto Target(M) | Auto PGA Target(M) | PGA Mem Alloc(M) | W/A PGA Used(M) | %PGA W/A Mem | %Auto W/A Mem | %Man W/A Mem | Global Mem Bound(K) |
|---|--------------------|--------------------|------------------|-----------------|--------------|---------------|--------------|---------------------|
| B | 20,480 | 11,154 | 11,174.79 | 400.92 | 3.59 | 100.00 | 0.00 | 1,048,576 |
| E | 20,480 | 9,640 | 16,496.97 | 3,108.22 | 18.84 | 100.00 | 0.00 | 1,048,576 |

[Back to Advisory Statistics](#)

[Back to Top](#)

PGA Aggr Target Histogram

- Optimal Executions are purely in-memory operations

| Low Optimal | High Optimal | Total Execs | Optimal Execs | 1-Pass Execs | M-Pass Execs |
|-------------|--------------|-------------|---------------|--------------|--------------|
| 2K | 4K | 492,432 | 492,432 | 0 | 0 |
| 64K | 128K | 9,833 | 9,833 | 0 | 0 |
| 128K | 256K | 287 | 287 | 0 | 0 |
| 256K | 512K | 156 | 156 | 0 | 0 |
| 512K | 1024K | 854 | 854 | 0 | 0 |
| 1M | 2M | 1,240 | 1,240 | 0 | 0 |
| 2M | 4M | 516 | 516 | 0 | 0 |
| 4M | 8M | 283 | 268 | 15 | 0 |

PGA Memory Advisory

pga_aggregate_target过小会导致PGA overalloc过载，但对于变态的HASH/SORT需求，再大的PGA也达不到cache hit 100%

PGA Memory Advisory

- When using Auto Memory Mgmt, minimally choose a pga_aggregate_target value where Estd PGA Overalloc Count is 0

| PGA Target Est (MB) | Size Factr | W/A MB Processed | Estd Extra W/A MB Read/ Written to Disk | Estd PGA Cache Hit % | Estd PGA Overalloc Count | Estd Time |
|---------------------|------------|------------------|---|----------------------|--------------------------|---------------|
| 2,560 | 0.13 | 10,466,791.54 | 19,759,279.75 | 35.00 | 105,836 | 8,713,283,754 |
| 5,120 | 0.25 | 10,466,791.54 | 19,446,281.02 | 35.00 | 104,802 | 8,623,055,463 |
| 10,240 | 0.50 | 10,466,791.54 | 2,736,989.43 | 79.00 | 2,099 | 3,806,260,135 |
| 15,360 | 0.75 | 10,466,791.54 | 2,688,456.12 | 80.00 | 1,953 | 3,792,269,415 |
| 20,480 | 1.00 | 10,466,791.54 | 1,449,037.51 | 88.00 | 1,128 | 3,434,981,631 |
| 24,576 | 1.20 | 10,466,791.54 | 1,383,030.93 | 88.00 | 310 | 3,415,953,882 |
| 28,672 | 1.40 | 10,466,791.54 | 1,383,030.93 | 88.00 | 0 | 3,415,953,882 |
| 32,768 | 1.60 | 10,466,791.54 | 1,382,883.73 | 88.00 | 0 | 3,415,911,450 |
| 36,864 | 1.80 | 10,466,791.54 | 1,382,883.73 | 88.00 | 0 | 3,415,911,450 |
| 40,960 | 2.00 | 10,466,791.54 | 1,382,883.73 | 88.00 | 0 | 3,415,911,450 |
| 61,440 | 3.00 | 10,466,791.54 | 1,382,883.73 | 88.00 | 0 | 3,415,911,450 |
| 81,920 | 4.00 | 10,466,791.54 | 1,382,883.73 | 88.00 | 0 | 3,415,911,450 |
| 122,880 | 6.00 | 10,466,791.54 | 1,382,883.73 | 88.00 | 0 | 3,415,911,450 |
| 163,840 | 8.00 | 10,466,791.54 | 1,382,883.73 | 88.00 | 0 | 3,415,911,450 |

Load Profile- Logons 登陆

Logons 登陆次数, **logon storm** 登陆风暴, 结合**AUDIT**审计数据一起看。
短连接的附带效应是游标缓存无用, 以下为短连接变长连接后的优化效果

变更前

| | Per Second | Per Transaction |
|------------------|------------|-----------------|
| Redo size: | 244,606.59 | 13,269.94 |
| Logical reads: | 5,964.59 | 323.58 |
| Block changes: | 1,278.41 | 69.35 |
| Physical reads: | 339.03 | 18.39 |
| Physical writes: | 35.30 | 1.92 |
| User calls: | 693.44 | 37.62 |
| Parses: | 241.46 | 13.10 |
| Hard parses: | 0.16 | 0.01 |
| Sorts: | 97.93 | 5.31 |
| Logons: | 16.05 | 0.87 |
| Executes: | 617.55 | 33.50 |
| Transactions: | 18.43 | |

| Statistic Name | Time (s) | % of DB Time |
|--|----------|--------------|
| sql execute elapsed time | 2,603.73 | 73.47 |
| DB CPU | 2,430.37 | 68.58 |
| connection management call elapsed time | 511.90 | 14.45 |
| parse time elapsed | 163.60 | 4.62 |
| PL/SQL execution elapsed time | 84.88 | 2.40 |
| hard parse elapsed time | 27.08 | 0.76 |
| sequence load elapsed time | 17.88 | 0.50 |
| hard parse (sharing criteria) elapsed time | 0.01 | 0.00 |
| repeated bind elapsed time | 0.00 | 0.00 |
| DB time | 3,543.74 | |
| background elapsed time | 513.68 | |
| background cpu time | 351.72 | |

变更后

| | Per Second | Per Transaction |
|------------------|------------|-----------------|
| Redo size: | 314,037.68 | 4,249.08 |
| Logical reads: | 7,939.19 | 107.42 |
| Block changes: | 1,629.35 | 22.05 |
| Physical reads: | 221.23 | 2.99 |
| Physical writes: | 41.85 | 0.57 |
| User calls: | 1,005.17 | 13.60 |
| Parses: | 76.15 | 1.03 |
| Hard parses: | 0.16 | 0.00 |
| Sorts: | 37.36 | 0.51 |
| Logons: | 0.36 | 0.00 |
| Executes: | 810.16 | 10.96 |
| Transactions: | 73.91 | |

| Statistic Name | Time (s) | % of DB Time |
|--|----------|--------------|
| DB CPU | 1,661.42 | 74.02 |
| sql execute elapsed time | 1,558.64 | 69.44 |
| PL/SQL execution elapsed time | 66.66 | 2.97 |
| parse time elapsed | 37.24 | 1.66 |
| hard parse elapsed time | 15.09 | 0.67 |
| connection management call elapsed time | 8.37 | 0.37 |
| sequence load elapsed time | 3.53 | 0.16 |
| PL/SQL compilation elapsed time | 0.49 | 0.02 |
| hard parse (sharing criteria) elapsed time | 0.08 | 0.00 |
| failed parse elapsed time | 0.08 | 0.00 |
| repeated bind elapsed time | 0.00 | 0.00 |
| DB time | 2,244.66 | |
| background elapsed time | 669.28 | |
| background cpu time | 382.82 | |

Load Profile-Execute、rollback、transactions

Executes 执行次数，反应执行频率

Rollback 回滚次数，反应回滚频率，但是这个指标不太精确，参考而已，别太当真

Transactions 每秒事务数，和数据库层的TPS，可以看做压力测试或比对性能是的一个指标，孤立看无意义。

| | | | |
|-------------------|----------|-----|--|
| W/A MB processed: | 26.1 | 0.0 | |
| Logons: | 0.9 | 0.0 | |
| Executes: | 15,591.8 | 8.4 | |
| Rollbacks: | 100.5 | 0.1 | |
| Transactions: | 1,861.2 | | |

Instance Efficiency Percentages (Target 100%)

基于命中率的调优方法论已经过时，但仍具有参考价值

Instance Efficiency Percentages (Target 100%)

| | | | |
|-------------------------------|-------|-------------------|--------|
| Buffer Nowait %: | 99.99 | Redo NoWait %: | 99.98 |
| Buffer Hit %: | 98.46 | In-memory Sort %: | 100.00 |
| Library Hit %: | 95.54 | Soft Parse %: | 67.05 |
| Execute to Parse %: | 90.43 | Latch Hit %: | 91.46 |
| Parse CPU to Parse Elapsed %: | 2.34 | % Non-Parse CPU: | 93.62 |

全部是越高越好！

Buffer nowait%: session 申请一个buffer(兼容模式)不等待的次数比例。

Buffer HIT%: 经典的经典，高速缓存命中率，反应物理读和缓存命中间的纠结，但这个指标即便99% 也不能说明物理读等待少了

Redo nowait%: session 在生成redo entry时不用等待的比例，redo相关的资源争用例如redo space request争用可能造成生成redo时需求等待。此项数据来源于v\$sysstat中的redo log space requests/redu entries

Instance Efficiency Percentages (Target 100%) In-Memory Sort%

基于命中率的调优方法论已经过时，但仍具有参考价值

Instance Efficiency Percentages (Target 100%)

| | | | |
|-------------------------------|-------|-------------------|--------|
| Buffer Nowait %: | 99.99 | Redo NoWait %: | 99.98 |
| Buffer Hit %: | 98.46 | In-memory Sort %: | 100.00 |
| Library Hit %: | 95.54 | Soft Parse %: | 67.05 |
| Execute to Parse %: | 90.43 | Latch Hit %: | 91.46 |
| Parse CPU to Parse Elapsed %: | 2.34 | % Non-Parse CPU: | 93.62 |

In-memory Sort%:这个指标因为它不计算workarea中所有的操作类型，所以现在越来越鸡肋了。纯粹在内存中完成的排序比例。数据来源于v\$sqlstat statistics sorts (disk) 和 sorts (memory).

| | | | |
|----------------|-------------|------------|--------|
| sorts (disk) | 45 | 0.01 | 0.00 |
| sorts (memory) | 1,154,835 | 369.76 | 0.20 |
| sorts (rows) | 911,885,722 | 291,973.66 | 156.87 |

Instance Efficiency Percentages (Target 100%)

基于命中率的调优方法论已经过时，但仍具有参考价值

Instance Efficiency Percentages (Target 100%)

| | | | |
|-------------------------------|-------|-------------------|--------|
| Buffer Nowait %: | 99.99 | Redo NoWait %: | 99.98 |
| Buffer Hit %: | 98.46 | In-memory Sort %: | 100.00 |
| Library Hit %: | 95.54 | Soft Parse %: | 67.05 |
| Execute to Parse %: | 90.43 | Latch Hit %: | 91.46 |
| Parse CPU to Parse Elapsed %: | 2.34 | % Non-Parse CPU: | 93.62 |

Library Hit%: library cache命中率，申请一个library cache object例如一个SQL cursor时，其已经在library cache中的比例。数据来源 V\$librarycache的pins和pinhits。合理值：>95%

Soft Parse: 软解析比例，无需多说的经典指标，数据来源v\$sysstat statistics的parse count(total)和parse count(hard)。合理值>95%

| | | | |
|---------------------------|------------|-----------|------|
| opened cursors cumulative | 27,000,000 | 0,000,000 | 0.00 |
| parse count (describe) | 72 | 0.02 | 0.00 |
| parse count (failures) | 1,830 | 0.59 | 0.00 |
| parse count (hard) | 1,536,168 | 491.86 | 0.26 |
| parse count (total) | 4,661,978 | 1,492.70 | 0.80 |

Instance Efficiency Percentages (Target 100%)

基于命中率的调优方法论已经过时，但仍具有参考价值

Instance Efficiency Percentages (Target 100%)

| | | | |
|-------------------------------|-------|-------------------|--------|
| Buffer Nowait %: | 99.99 | Redo NoWait %: | 99.98 |
| Buffer Hit %: | 98.46 | In-memory Sort %: | 100.00 |
| Library Hit %: | 95.54 | Soft Parse %: | 67.05 |
| Execute to Parse %: | 90.43 | Latch Hit %: | 91.46 |
| Parse CPU to Parse Elapsed %: | 2.34 | % Non-Parse CPU: | 93.62 |

Execute to Parse% 指标反映了执行解析比 其公式为 $1 - (\text{parse}/\text{execute})$ ，目标为100% 及接近于只 执行而不解析。 数据来源 `v$sysstat statistics parse count (total)` 和 `execute count`

Latch Hit%: willing-to-wait latch 申请不要等待的比例。 数据来源 `V$latch gets` 和 `misses`

Parse CPU To Parse Elapsed: 该指标反映了 快照内解析CPU时间和总的解析时间的比值 ($\text{Parse CPU Time} / \text{Parse Elapsed Time}$)； 若该指标水平很低，那么说明在整个解析过程中 实际在CPU上运算的时间是很短的，而主要的解析时间都耗费在各种其他非空闲的等待事件上了(如 `latch:shared pool`, `row cache lock` 之类等) 数据来源 `V$sysstat` 的 `parse time cpu` 和 `parse time elapsed`

%Non-Parse CPU 非解析cpu比例，公式为 $(\text{DB CPU} - \text{Parse CPU}) / \text{DB CPU}$ ， 若大多数CPU都用在解析上了，则可能好钢没用在刀上了。 数据来源 `v$sysstat` 的 `parse time cpu` 和 `cpu used by this session`

Shared Pool Statistics

反应SQL重用率和共享池中cursor对内存的使用

Shared Pool Statistics

| | Begin | End |
|----------------------------|-------|-------|
| Memory Usage %: | 96.53 | 98.49 |
| % SQL with executions>1: | 84.79 | 65.92 |
| % Memory for SQL w/exec>1: | 80.54 | 53.02 |

该环节提供一个大致的SQL重用及shared pool内存使用的评估。应用是否共享SQL? 有多少内存是给只运行一次的SQL占掉的, 对比共享SQL呢?

如果该环节中% SQL with executions>1的比例 小于%90, 考虑用下面链接的SQL去抓 硬编码的非绑定变量SQL语句。

利用FORCE_MATCHING_SIGNATURE捕获非绑定变量SQL

http://www.askmaclean.com/archives/%E5%88%A9%E7%94%A8force_matching_signature%E6%8D%95%E8%8E%B7%E9%9D%9E%E7%BB%91%E5%AE%9A%E5%8F%98%E9%87%8Fsql.html

Top 5 Timed Foreground Events → 舞会主角

Top 5 万众瞩目，DBA为你倾倒！

Top 5 Timed Foreground Events

| Event | Waits | Time(s) | Avg wait (ms) | % DB time | Wait Class |
|---------------------------------|-----------|---------|---------------|-----------|------------|
| DB CPU | | 14,703 | | 30.77 | |
| log file sync | 723,365 | 12,898 | 18 | 27.00 | Commit |
| gc buffer busy acquire | 313,547 | 5,390 | 17 | 11.28 | Cluster |
| gc current block busy | 258,139 | 5,385 | 21 | 11.27 | Cluster |
| cell single block physical read | 4,525,253 | 5,125 | 1 | 10.73 | User I/O |

基于Wait Interface的调优是目前的主流！每个指标都重要！

基于命中比例的调优，好比是统计局的报告，张财主家财产100万，李木匠家财产1万，平均财产50.5万。

基于等待事件的调优，好比马路上100辆汽车的行驶记录表，上车用了几分钟，红灯等了几分钟，拥堵塞了几分钟。。。

Mysql梦寐以求的东西.....

Top 5 Timed Foreground Events DB CPU/Cpu Time

Top 5 万众瞩目，DBA为你倾倒！

Top 5 Timed Foreground Events

| Event | Waits | Time(s) | Avg wait (ms) | % DB time | Wait Class |
|---------------------------------|-----------|---------|---------------|-----------|------------|
| DB CPU | | 14,703 | | 30.77 | |
| log file sync | 723,365 | 12,898 | 18 | 27.00 | Commit |
| gc buffer busy acquire | 313,547 | 5,390 | 17 | 11.28 | Cluster |
| gc current block busy | 258,139 | 5,385 | 21 | 11.27 | Cluster |
| cell single block physical read | 4,525,253 | 5,125 | 1 | 10.73 | User I/O |

CPU 上在干什么？

逻辑读？ 解析？ Latch spin? PL/SQL、函数运算？

DB CPU/CPU time是Top 1 是好事情吗？ 未必！

注意DB CPU不包含 wait on cpu queue！

Top 5 Timed Foreground Events DB CPU/Cpu Time

Top 5 万众瞩目，DBA为你倾倒！

结合Host CPU、Instance CPU、SQL ordered by CPU Time一起看哦！

Host CPU (CPUs: 16 Cores: 8 Sockets: 2)

| Load Average Begin | Load Average End | %User | %System | %WIO | %Idle |
|--------------------|------------------|-------|---------|------|-------|
| 5.10 | 7.43 | 26.9 | 3.0 | 0.0 | 67.7 |

Instance CPU

| %Total CPU | %Busy CPU | %DB time waiting for CPU (Resource Manager) |
|------------|-----------|---|
| 30.4 | 94.3 | 0.0 |

SQL ordered by CPU Time

- Resources reported for PL/SQL code includes the resources used by all SQL statements called by the code.
- %Total - CPU Time as a percentage of Total DB CPU
- %CPU - CPU Time as a percentage of Elapsed Time
- %IO - User I/O Time as a percentage of Elapsed Time
- Captured SQL account for 19.1% of Total CPU Time (s): 14,703
- Captured PL/SQL account for 0.8% of Total CPU Time (s): 14,703

| CPU Time (s) | Executions | CPU per Exec (s) | %Total | Elapsed Time (s) | %CPU | %IO | SQL Id | SQL Module | SQL Text |
|--------------|------------|------------------|--------|------------------|-------|-------|---------------|---------------------------------|-----------------------------------|
| 378.93 | 442,560 | 0.00 | 2.58 | 1,612.55 | 23.50 | 65.01 | d36p2zk77mu80 | dm_oracle@aaa-brm02 (TNS V1-V3) | select distinct poid_DB, poid_... |
| 272.63 | 3,264,486 | 0.00 | 1.85 | 922.66 | 29.55 | 50.71 | 61rcbymrd925x | dm_oracle@aaa-brm02 (TNS V1-V3) | select poid_DB, poid_ID0, poid... |
| 242.10 | 467,311 | 0.00 | 1.65 | 260.80 | 92.83 | 0.00 | 56nmjy6g1aky8 | dm_oracle@aaa-brm05 (TNS V1-V3) | select distinct poid_DB, poid_... |
| 193.67 | 112,084 | 0.00 | 1.32 | 204.32 | 94.79 | 0.00 | 2ckhwymh5d72g | JDBC Thin Client | SELECT B.SERVICEID AS PRODUCT... |
| 150.15 | 442,554 | 0.00 | 1.02 | 302.68 | 49.61 | 37.11 | 3q72620j7sw4q | dm_oracle@aaa-brm05 (TNS V1-V3) | select distinct poid_DB, poid_... |
| 122.77 | 112,086 | 0.00 | 0.84 | 133.03 | 92.29 | 34.84 | cp8azjyvnn9uv | JDBC Thin Client | BEGIN p_diy_partnersub_qr (:1... |
| 121.37 | 3,367,715 | 0.00 | 0.83 | 134.94 | 89.94 | 0.00 | adc95dp9cgt6h | dm_oracle@aaa-brm02 (TNS V1-V3) | select poid_DB, poid_ID0, poid... |

Operating System Statistics

- *TIME statistic values are diffed. All others display actual values. End Value is displayed if different
- ordered by statistic type (CPU Use, Virtual Memory, Hardware Config), Name

| Statistic | Value | End Value |
|-----------------------|----------------|-----------|
| BUSY_TIME | 1,709,372 | |
| IDLE_TIME | 3,585,824 | |
| IOWAIT_TIME | 360 | |
| NICE_TIME | 0 | |
| SYS_TIME | 157,135 | |
| USER_TIME | 1,424,443 | |
| LOAD | 5 | 7 |
| PHYSICAL_MEMORY_BYTES | 75,804,418,048 | |
| NUM_CPUS | 16 | |
| NUM_CPU_CORES | 8 | |
| NUM_CPU_SOCKETS | 2 | |

db file sequential read- Top event

Avg wait time应当小于20ms

”db file sequential read”单块读等待是一种最为常见的物理IO等待事件，这里的sequential指的是将数据块读入到相连的内存空间中(contiguous memory space)，而不是指所读取的数据块是连续的。该wait event可能在以下情景中发生：

<http://www.askmaclean.com/archives/db-file-sequential-read-wait-event.html>

| | | | |
|------------------------------|---------------|------------|----------|
| sql area evicted | 2 | 0.00 | 0.00 |
| sql area purged | 1,494 | 0.41 | 0.00 |
| summed dirty queue length | 16,125,464 | 4,470.60 | 23.11 |
| switch current to new buffer | 115,043 | 31.89 | 0.16 |
| table fetch by rowid | 122,174,123 | 33,871.38 | 175.11 |
| table fetch continued row | 483,100 | 133.93 | 0.69 |
| table scan blocks gotten | 60,861,860 | 16,873.25 | 87.23 |
| table scan rows gotten | 1,097,387,565 | 304,238.14 | 1,572.89 |
| table scans (direct read) | 0 | 0.00 | 0.00 |
| table scans (long tables) | 0 | 0.00 | 0.00 |
| table scans (short tables) | 1,161,933 | 322.13 | 1.67 |
| temp space allocated (bytes) | 0 | 0.00 | 0.00 |

db file scattered read - Top event

Avg wait time应当小于20ms

常见原因 Fast Full scan Index , FULL SCAN large table

| | | | |
|---|---------------|------------|----------|
| immediate (CR) block cleanout applications | 921,121 | 92.91 | 0.20 |
| immediate (CURRENT) block cleanout applications | 1,535,887 | 98.34 | 0.53 |
| index crx upgrade (positioned) | 8,306 | 0.53 | 0.00 |
| index crx upgrade (prefetch) | 20,551 | 1.32 | 0.01 |
| index fast full scans (full) | 4 | 0.00 | 0.00 |
| index fetch by key | 5,708,408,652 | 365,483.83 | 1,958.45 |
| index scans kdiixs1 | 3,866,359,624 | 247,545.69 | 1,326.48 |
| leaf node 90-10 splits | 3,425 | 0.22 | 0.00 |
| leaf node splits | 125,845 | 8.06 | 0.04 |

| | | | |
|-----------------------------------|----------------|--------------|-----------|
| switch current to new buffer | 1,264,886 | 80.98 | 0.43 |
| table fetch by rowid | 14,382,569,915 | 920,851.52 | 4,934.40 |
| table fetch continued row | 699,309,025 | 44,773.62 | 239.92 |
| table scan blocks gotten | 949,871,358 | 60,816.01 | 325.88 |
| table scan rows gotten | 30,300,246,632 | 1,939,989.05 | 10,395.47 |
| table scans (direct read) | 6,549 | 0.42 | 0.00 |
| table scans (long tables) | 6,646 | 0.43 | 0.00 |
| table scans (short tables) | 1,317,722 | 84.37 | 0.45 |
| total cf enq hold time | 23,019 | 1.47 | 0.01 |
| total number of cf enq holders | 1,880 | 0.12 | 0.00 |
| total number of slots | 14 | 0.00 | 0.00 |
| total number of times SMON posted | 3,549 | 0.23 | 0.00 |
| transaction lock background gets | 0 | 0.00 | 0.00 |

db file sequential/scattered read - Top event

SQL ordered by Reads

- %Total - Physical Reads as a percentage of Total Disk Reads
- %CPU - CPU Time as a percentage of Elapsed Time
- %IO - User I/O Time as a percentage of Elapsed Time
- Total Disk Reads: 65,116,233
- Captured SQL account for 22.5% of Total

| Physical Reads | Executions | Reads per Exec | % Total | Elapsed Time (s) | %CPU | %IO | SQL Id |
|----------------|------------|----------------|---------|------------------|-------|-------|-------------------------------|
| 2,923,909 | 18,869 | 154.96 | 4.49 | 13,775.98 | 6.97 | 88.61 | gf37z0zzxc450 |
| 1,116,277 | 28 | 39,867.04 | 1.71 | 2,335.72 | 48.32 | 49.64 | bcd3f84dbcv08 |
| 959,749 | 22 | 43,624.95 | 1.47 | 278.63 | 25.73 | 54.41 | 3tzqjzqg0s45v |
| 674,218 | 123 | 5,481.45 | 1.04 | 6,854.08 | 33.86 | 61.15 | 5v5na7fj9g1dh |
| 595,396 | 763 | 780.34 | 0.91 | 16,164.52 | 30.72 | 33.63 | 2j8j6sw668r3t |
| 379,383 | 2,634 | 144.03 | 0.58 | 6,330.32 | 0.87 | 0.57 | 9k6vh8jb5t7pt |
| 348,368 | 161 | 2,163.78 | 0.53 | 2,298.74 | 4.59 | 91.68 | 9vcxgb8ymva9p |
| 345,962 | 177 | 1,954.59 | 0.53 | 2,341.43 | 5.10 | 91.38 | b6n6jbdp1rtn0 |
| 344,561 | 4,691 | 73.45 | 0.53 | 1,221.24 | 17.59 | 73.75 | 5s8jd5fd611cf |
| 312,689 | 8 | 39,086.13 | 0.48 | 447.88 | 43.99 | 55.85 | f35w06f7k3868 |


Segments by Physical Reads

- Total Physical Reads: 3,610,975
- Captured Segments account for 78.7% of Total

| Owner | Tablespace Name | Object Name | Subobject Name | Obj. Type | Physical Reads | % Total |
|-------|-----------------|-------------------------|----------------|-----------|----------------|---------|
| PIN | PIN02 | PURCHASED_PRODUCT_T | | TABLE | 650,031 | 18.00 |
| PIN | PINX02 | I_PURCHASED_PRODUCT__ID | | INDEX | 459,291 | 12.72 |
| PIN | PIN01 | ACCOUNT_T | | TABLE | 171,294 | 4.74 |
| PIN | PINX00 | I_SERVICE__LOGIN | | INDEX | 170,862 | 4.73 |
| PIN | PIN02 | SERVICE_T | | TABLE | 168,201 | 4.66 |

Log file sync蝴蝶效应

Log file sync → enq: TX , gc buffer busy, buffer busy wait
等待事件的混沌理论，性能不是线性的，而是多纬度的



Log file sync实战案例

Log file sync → enq: TX , gc buffer busy, buffer busy wait , enq:TX
index contention

等待事件的混沌理论，性能不是线性的，而是多纬度的

Top 5 Timed Foreground Events

| Event | Waits | Time(s) | Avg wait (ms) | % DB time | Wait Class |
|---------------------------------|-----------|---------|---------------|-----------|-------------|
| enq: TX - row lock contention | 170,732 | 76,510 | 448 | 41.21 | Application |
| gc buffer busy release | 311,756 | 26,611 | 85 | 14.33 | Cluster |
| gc buffer busy acquire | 598,333 | 20,626 | 34 | 11.11 | Cluster |
| cell single block physical read | 3,588,957 | 14,903 | 4 | 8.03 | User I/O |
| log file sync | 342,534 | 13,630 | 40 | 7.34 | Commit |

Global Cache and Enqueue Services - Workload Characteristics

| | |
|---|------|
| Avg global enqueue get time (ms): | 41.0 |
| Avg global cache cr block receive time (ms): | 6.0 |
| Avg global cache current block receive time (ms): | 8.9 |
| Avg global cache cr block build time (ms): | 0.0 |
| Avg global cache cr block send time (ms): | 0.0 |
| Global cache log flushes for cr blocks served %: | 9.7 |
| Avg global cache cr block flush time (ms): | 34.8 |
| Avg global cache current block pin time (ms): | 8.4 |
| Avg global cache current block send time (ms): | 0.0 |
| Global cache log flushes for current blocks served %: | 9.9 |
| Avg global cache current block flush time (ms): | 37.4 |

Log file sync实战案例

log file parallel write慢=> log file sync慢=>commit慢, commit慢则释放行锁慢。Rac flush redo也受到写redo慢的影响, 则出现gc buffer busy release/acquire, 前后相互作用→ enq:TX 大幅出现

Top 5 Timed Foreground Events

| Event | Waits | Time(s) | Avg wait (ms) | % DB time | Wait Class |
|---------------------------------|-----------|---------|---------------|-----------|-------------|
| enq: TX - row lock contention | 170,732 | 76,510 | 448 | 41.21 | Application |
| gc buffer busy release | 311,756 | 26,611 | 85 | 14.33 | Cluster |
| gc buffer busy acquire | 598,333 | 20,626 | 34 | 11.11 | Cluster |
| cell single block physical read | 3,588,957 | 14,903 | 4 | 8.03 | User I/O |
| log file sync | 342,534 | 13,630 | 40 | 7.34 | Commit |

| | | | |
|----------------|------------|-----------|--------|
| user calls | 41,091,880 | 11,437.02 | 120.22 |
| user commits | 340,894 | 94.88 | 1.00 |
| user rollbacks | 902 | 0.25 | 0.00 |

| Event | Waits | % Time -outs | Total Wait Time (s) | Avg wait (ms) | Waits /txn | % bg time |
|-----------------------------|-----------|--------------|---------------------|---------------|------------|-----------|
| db file parallel write | 488,325 | 0 | 4,231 | 9 | 1.43 | 35.8 |
| acs log flush sync | 6,828,951 | 0 | 3,648 | 1 | 19.98 | 30.9 |
| log file parallel write | 167,338 | 0 | 2,909 | 17 | 0.49 | 24.6 |
| log file sequential read | 468 | 0 | 86 | 183 | 0.00 | 0.7 |
| control file parallel write | 1,765 | 0 | 8.4 | 48 | 0.04 | 0.7 |

Global Cache and Enqueue Services - Workload Characteristics

| | |
|---|------|
| Avg global enqueue get time (ms): | 41.0 |
| Avg global cache cr block receive time (ms): | 6.0 |
| Avg global cache current block receive time (ms): | 8.9 |
| Avg global cache cr block build time (ms): | 0.0 |
| Avg global cache cr block send time (ms): | 0.0 |
| Global cache log flushes for cr blocks served %: | 9.7 |
| Avg global cache cr block flush time (ms): | 34.8 |
| Avg global cache current block pin time (ms): | 8.4 |
| Avg global cache current block send time (ms): | 0.0 |
| Global cache log flushes for current blocks served %: | 9.9 |
| Avg global cache current block flush time (ms): | 37.4 |

Log file sync实战案例

Enq:TX row lock出现在哪里？哪些语句受到GC buffer busy影响？

最主要是update和insert受影响，前台处理业务速度放慢。方向对了就处处对得上了

SQL ordered by Cluster Wait Time

- %Total - Cluster Time as a percentage of Total Cluster Wait Time
- %Clu - Cluster Time as a percentage of Elapsed Time
- %CPU - CPU Time as a percentage of Elapsed Time
- %IO - User I/O Time as a percentage of Elapsed Time
- Only SQL with Cluster Wait Time > .005 seconds is reported
- Total Cluster Wait Time (s): 67,176
- Captured SQL account for 93.8% of Total

| Cluster Wait Time (s) | Executions | %Total | Elapsed Time(s) | %Clu | %CPU | %IO | SQL Id | SQL Module | SQL Text |
|-----------------------|------------|--------|-----------------|-------|------|-------|----------------|---------------------------------|-----------------------------------|
| 23,377.70 | 27,253 | 23.86 | 97,970.74 | 34.80 | 0.04 | 0.00 | 4636378wmh2bp | dm_oracle@aaa-brm02 (TNS V1-V3) | update data_t set poid_rev = p... |
| 15,998.59 | 94,272 | 84.29 | 18,980.50 | 23.82 | 0.54 | 1.74 | q71uiftrcn0npd | dm_oracle@aaa-brm02 (TNS V1-V3) | insert into event_t (poid_DB,... |
| 8,743.97 | 57,401 | 87.68 | 9,972.85 | 13.02 | 0.58 | 3.48 | 0v25hwd7t2bus | dm_oracle@aaa-brm02 (TNS V1-V3) | insert into billog_t (poid_D... |
| 5,881.08 | 41,794 | 96.20 | 6,113.09 | 8.75 | 0.51 | 0.59 | 87vd326hc3ar0 | dm_oracle@aaa-brm05 (TNS V1-V3) | insert into event_t (poid_DB,... |
| 1,634.56 | 27,255 | 77.46 | 2,110.12 | 2.43 | 1.30 | 21.21 | 5c0d20xn9qbhm | dm_oracle@aaa-brm05 (TNS V1-V3) | insert into event_t (poid_DB,... |
| 757.75 | 12,214 | 93.24 | 812.67 | 1.13 | 0.90 | 0.05 | 1f9918t43pbmq | dm_oracle@aaa-brm05 (TNS V1-V3) | insert into using_trafficaccou... |
| 623.05 | 27,800 | 93.04 | 669.65 | 0.93 | 2.80 | 1.47 | 5iicibv7bqz5g | dm_oracle@aaa-brm05 (TNS V1-V3) | delete from event_t where poid... |
| 599.40 | 444,105 | 41.22 | 1,454.16 | 0.89 | 5.70 | 53.87 | 6w4q4kh8cprmu | dm_oracle@aaa-brm02 (TNS V1-V3) | select poid_DB, poid_ID0, poid... |
| 553.35 | 12,261 | 99.12 | 558.24 | 0.82 | 0.71 | 0.19 | awas7if1vmrdq | dm_oracle@aaa-brm02 (TNS V1-V3) | update service_t set poid_rev ... |
| 544.39 | 2,639,714 | 22.49 | 2,420.87 | 0.81 | 8.97 | 68.87 | 61rcbymrd925x | dm_oracle@aaa-brm02 (TNS V1-V3) | select poid_DB, poid_ID0, poid... |
| 471.91 | 12,261 | 97.33 | 484.85 | 0.70 | 1.58 | 1.17 | 302qc7nqv9d9 | dm_oracle@aaa-brm02 (TNS V1-V3) | insert into service_t (poid_D... |
| 402.94 | 12,263 | 96.07 | 419.42 | 0.60 | 1.73 | 1.68 | 7n20v2q57mxsg | dm_oracle@aaa-brm05 (TNS V1-V3) | insert into account_t (poid_D... |
| 372.29 | 396,814 | 9.48 | 3,926.10 | 0.55 | 8.50 | 82.58 | d36p2zk77mu80 | dm_oracle@aaa-brm02 (TNS V1-V3) | select distinct poid_DB, poid... |
| 358.30 | 444,107 | 29.82 | 1,201.71 | 0.53 | 6.22 | 64.93 | 607hhqizab9h6 | dm_oracle@aaa-brm02 (TNS V1-V3) | select state from account_name... |
| 305.18 | 60,000 | 44.40 | 744.57 | 0.45 | 1.74 | 54.00 | 5mwa5mwa5mwa5 | dm_oracle@aaa-brm02 (TNS V1-V3) | insert into ... |

Segments by Row Lock Waits

- % of Capture shows % of row lock waits for each top segment compared
- with total row lock waits for all segments captured by the Snapshot

| Owner | Tablespace Name | Object Name | Subobject Name | Obj. Type | Row Lock Waits | % of Capture |
|-------|-----------------|-----------------------|----------------|-----------------|----------------|--------------|
| PIN | PIN02 | DATA_T | | TABLE | 169,103 | 82.53 |
| PIN | PINX01 | I_EVENT_EVENT_NO | P_R_02292012 | INDEX PARTITION | 15,273 | 7.45 |
| PIN | PINX00 | IDX_BILLLOG_TIMESTAMP | P_R_02292012 | INDEX PARTITION | 9,321 | 4.55 |
| PIN | PINX00 | BILL_LOG_IDX | P_R_02292012 | INDEX PARTITION | 3,929 | 1.92 |
| PIN | PINX02 | I_EVENT_SERVOBJ_END_T | P_R_02292012 | INDEX PARTITION | 2,481 | 1.21 |

Log file sync实战案例

Global Buffer Busy → gc buffer busy acquire/release 受影响的segment

Segments by Global Cache Buffer Busy

- % of Capture shows % of GC Buffer Busy for each top segment compared
- with GC Buffer Busy for all segments captured by the Snapshot


| Owner | Tablespace Name | Object Name | Subobject Name | Obj. Type | GC Buffer Busy | % of Capture |
|-------|-----------------|-----------------------|----------------|-----------------|----------------|--------------|
| PIN | PIN00 | EVENT_T | P_R_02292012 | TABLE PARTITION | 183,367 | 24.90 |
| PIN | PINX01 | I_EVENT_EVENT__NO | P_R_02292012 | INDEX PARTITION | 161,690 | 21.96 |
| PIN | PIN02 | DATA_T | | TABLE | 157,531 | 21.39 |
| PIN | PINX00 | IDX_BILLLOG_TIMESTAMP | P_R_02292012 | INDEX PARTITION | 73,055 | 9.92 |
| PIN | PINX00 | BILL_LOG_IDX | P_R_02292012 | INDEX PARTITION | 67,141 | 9.12 |

| Wait Class | Waits | %Time -outs | Total Wait Time (s) | Avg wait (ms) | %DB time |
|---------------|------------|-------------|---------------------|---------------|----------|
| Application | 178,451 | 0 | 76,510 | 429 | 41.21 |
| Cluster | 6,357,009 | 0 | 67,164 | 11 | 36.17 |
| User I/O | 3,608,206 | 0 | 14,932 | 4 | 8.04 |
| Commit | 342,534 | 0 | 13,630 | 40 | 7.34 |
| DB CPU | | | 10,273 | | 5.53 |
| Concurrency | 258,669 | 0 | 3,658 | 14 | 1.97 |
| Other | 7,224,313 | 12 | 197 | 0 | 0.11 |
| Network | 41,094,256 | 0 | 47 | 0 | 0.03 |
| Configuration | 252 | 11 | 31 | 121 | 0.02 |
| System I/O | 972 | 0 | 0 | 0 | 0.00 |

性能优化的多维度理论

- 增加了**cpu**→更大的并发量，更多的并发争用
- 调整了**io**存储→ 更少的**IO**，更多的**CPU**计算，更高的**cpu**使用率
- **Redo**写得慢→ 影响**commit**，造成**enq:tx**和**gc buffer busy**等待等
- **Datafile**写得慢→ 检查点完不成，日志无法切换，前台**DML hang**
- **Sequence nocache**→ **INSERT index**很容易造成**enq:index contention**，和**row cache lock**和 **enq:SQ**
- 通过数据库手段优化了性能→ 应用本身设计的瓶颈越来越凸显


不给应用开大手术，纯数据库优化的极限



经过Maclean
一番打造



但是即便再
打造也不可能。
。。。



才开了个头哦。。。。

**To Be
Continued.....**

つづく.....

敬请期待开Oracle调优鹰眼，深入理解AWR性能报告第二讲

更多信息

www.askmaclean.com



or

<http://www.askmaclean.com/archives/tag/tuning>

Question & Answer



If you have more questions later, feel free to ask